

weather_analysis

January 16, 2025

1 Global Weather Analysis

1.0.1 PM Accelerator Mission:

By making industry-leading tools and education available to individuals from all backgrounds, we level the playing field for future PM leaders. This is the PM Accelerator motto, as we grant aspiring and experienced PMs what they need most – Access. We introduce you to industry leaders, surround you with the right PM ecosystem, and discover the new world of AI product management skills.

1.1 Library Installation and Import

```
[1]: !pip install -r ../requirements.txt
```

```
[33]: import os
import zipfile
import csv
import string
import sys
import warnings
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
import random

import folium
import pandas as pd
from folium.plugins import HeatMap
from IPython.display import display

from xgboost import XGBRegressor
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import MinMaxScaler, StandardScaler, MaxAbsScaler,
↳ RobustScaler

from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
```

```

from sklearn.feature_selection import mutual_info_classif, f_classif, chi2,
    ↳SelectKBest
from sklearn.metrics import accuracy_score, precision_score, recall_score,
    ↳f1_score, classification_report, confusion_matrix, mean_squared_error,
    ↳mean_absolute_error
from sklearn.multiclass import OneVsRestClassifier as OVR
from sklearn.naive_bayes import MultinomialNB, ComplementNB, CategoricalNB
from sklearn.svm import LinearSVC

import tensorflow as tf
from keras.callbacks import ModelCheckpoint,
    ↳LearningRateScheduler, ReduceLROnPlateau, EarlyStopping
from keras.layers import Input, Dense, Dropout, BatchNormalization, Reshape,
    ↳Flatten, Bidirectional, LSTM, GRU
from keras.models import Model, Sequential
from keras.optimizers import Adam

```

time: 2.42 ms (started: 2025-01-16 03:17:43 +00:00)

```

[3]: SEED = 99
def random_seed(SEED):
    random.seed(SEED)
    os.environ['PYTHONHASHSEED'] = str(SEED)
    np.random.seed(SEED)
    tf.random.set_seed(SEED)
random_seed(SEED)
%load_ext autotime
%matplotlib inline

print(f"TensorFlow version: {tf.__version__}")
print("CUDA Version:", tf.sysconfig.get_build_info()["cuda_version"])
print("cuDNN Version:", tf.sysconfig.get_build_info()["cudnn_version"])

```

TensorFlow version: 2.15.0

CUDA Version: 12.2

cuDNN Version: 8

time: 4.61 ms (started: 2025-01-16 02:52:24 +00:00)

1.2 Data Cleaning & Preprocessing

```

[4]: df = pd.read_csv("/notebooks/data/GlobalWeatherRepository.csv")
df

```

```

[4]:
   country location_name latitude longitude  timezone \
0  Afghanistan      Kabul   34.5200    69.1800  Asia/Kabul
1   Albania      Tirana   41.3300    19.8200  Europe/Tirane
2   Algeria      Algiers   36.7600     3.0500  Africa/Algiers

```

3	Andorra	Andorra La Vella	42.5000	1.5200	Europe/Andorra
4	Angola	Luanda	-8.8400	13.2300	Africa/Luanda
...
47352	Venezuela	Caracas	10.5000	-66.9167	America/Caracas
47353	Vietnam	Hanoi	21.0333	105.8500	Asia/Bangkok
47354	Yemen	Sanaa	15.3547	44.2067	Asia/Aden
47355	Zambia	Lusaka	-15.4167	28.2833	Africa/Lusaka
47356	Zimbabwe	Harare	-17.8178	31.0447	Africa/Harare

	last_updated_epoch	last_updated	temperature_celsius	\
0	1715849100	2024-05-16 13:15	26.6	
1	1715849100	2024-05-16 10:45	19.0	
2	1715849100	2024-05-16 09:45	23.0	
3	1715849100	2024-05-16 10:45	6.3	
4	1715849100	2024-05-16 09:45	26.0	
...	
47352	1736934300	2025-01-15 05:45	15.5	
47353	1736936100	2025-01-15 17:15	23.1	
47354	1736937000	2025-01-15 13:30	18.9	
47355	1736937000	2025-01-15 12:30	28.0	
47356	1736937000	2025-01-15 12:30	25.6	

	temperature_fahrenheit	condition_text	...	air_quality_PM2.5	\
0	79.8	Partly Cloudy	...	8.400	
1	66.2	Partly cloudy	...	1.100	
2	73.4	Sunny	...	10.400	
3	43.3	Light drizzle	...	0.700	
4	78.8	Partly cloudy	...	183.400	
...	
47352	60.0	Light rain	...	2.590	
47353	73.6	Sunny	...	68.635	
47354	66.0	Sunny	...	26.455	
47355	82.4	Partly Cloudy	...	15.355	
47356	78.1	Patchy rain nearby	...	17.575	

	air_quality_PM10	air_quality_us-epa-index	air_quality_gb-defra-index	\
0	26.600	1	1	
1	2.000	1	1	
2	18.400	1	1	
3	0.900	1	1	
4	262.300	5	10	
...	
47352	2.960	1	1	
47353	69.005	4	9	
47354	69.745	2	3	
47355	15.355	1	2	
47356	17.760	2	2	

	sunrise	sunset	moonrise	moonset	moon_phase \
0	04:50 AM	06:50 PM	12:12 PM	01:11 AM	Waxing Gibbous
1	05:21 AM	07:54 PM	12:58 PM	02:14 AM	Waxing Gibbous
2	05:40 AM	07:50 PM	01:15 PM	02:14 AM	Waxing Gibbous
3	06:31 AM	09:11 PM	02:12 PM	03:31 AM	Waxing Gibbous
4	06:12 AM	05:55 PM	01:17 PM	12:38 AM	Waxing Gibbous
...
47352	06:50 AM	06:24 PM	08:11 PM	08:15 AM	Waning Gibbous
47353	06:36 AM	05:36 PM	06:57 PM	07:41 AM	Waning Gibbous
47354	06:33 AM	05:52 PM	07:23 PM	07:45 AM	Waning Gibbous
47355	05:48 AM	06:44 PM	08:13 PM	07:00 AM	Waning Gibbous
47356	05:33 AM	06:37 PM	08:05 PM	06:44 AM	Waning Gibbous

	moon_illumination
0	55
1	55
2	55
3	55
4	55
...	...
47352	98
47353	99
47354	99
47355	99
47356	99

[47357 rows x 41 columns]

time: 192 ms (started: 2025-01-16 02:52:24 +00:00)

1.2.1 Check Missing Values

```
[5]: missing_values=df.isnull().sum()
print("Missing values:\n", missing_values)
missing_values = df.isnull().sum().sum()
# Check if there are any missing values and print the result using print_
↪statement
if missing_values > 0:
    print(f"Missing values are present. Total missing values: {missing_values}")
else:
    print(f"No missing values are present in the Dataset.")
```

```
Missing values:
country          0
location_name    0
latitude         0
```

```

longitude      0
timezone       0
last_updated_epoch  0
last_updated    0
temperature_celsius  0
temperature_fahrenheit  0
condition_text   0
wind_mph        0
wind_kph        0
wind_degree     0
wind_direction  0
pressure_mb     0
pressure_in     0
precip_mm       0
precip_in       0
humidity        0
cloud           0
feels_like_celsius  0
feels_like_fahrenheit  0
visibility_km    0
visibility_miles  0
uv_index        0
gust_mph        0
gust_kph        0
air_quality_Carbon_Monoxide  0
air_quality_Ozone  0
air_quality_Nitrogen_dioxide  0
air_quality_Sulphur_dioxide  0
air_quality_PM2.5  0
air_quality_PM10  0
air_quality_us-epa-index  0
air_quality_gb-defra-index  0
sunrise         0
sunset          0
moonrise        0
moonset         0
moon_phase      0
moon_illumination  0
dtype: int64
No missing values are present in the Dataset.
time: 28.9 ms (started: 2025-01-16 02:52:24 +00:00)

```

1.2.2 Check Duplicated Values

```

[6]: duplicates_count = df.duplicated().sum()

# Check if there are any duplicate rows and print the result using f-strings

```

```

if df.duplicated().any():
    print(f"Duplicates are present. Total duplicate rows: {duplicates_count}")
else:
    print(f"No duplicates are present in the Dataset.")

```

No duplicates are present in the Dataset.
time: 80.4 ms (started: 2025-01-16 02:52:24 +00:00)

1.2.3 Check Infinity Values

```

[7]: # Check for inf values in numeric columns only
numeric_columns = df.select_dtypes(include=[np.number]).columns
inf_values_count = np.isinf(df[numeric_columns]).sum().sum()

# Check if there are any inf values and print the result
if inf_values_count > 0:
    print(f"Inf values are present. Total inf values: {inf_values_count}")
else:
    print(f"No inf values are present in the Dataset.")

```

No inf values are present in the Dataset.
time: 7.83 ms (started: 2025-01-16 02:52:24 +00:00)

1.2.4 Outlier Detection and Boxplot Visualization

This process involves detecting and removing outliers using the IQR (Interquartile Range) method, followed by boxplot visualization to understand the cleaned dataset.

- **Data Selection and Preparation**
 - Identifies numeric columns in the DataFrame using `select_dtypes()`
 - Creates an initial boolean mask for tracking outliers across all columns
- **Outlier Detection**
 - Calculates Q1 (25th percentile) and Q3 (75th percentile) for each numeric column
 - Computes IQR (Interquartile Range) and defines outlier boundaries using the $1.5 * \text{IQR}$ rule
 - Marks data points outside these boundaries as outliers in the mask
- **Data Cleaning**
 - Creates a new DataFrame `df_no_outliers` by filtering out the identified outliers using the mask
- **Visualization Setup**
 - Calculates the required number of rows and columns for subplot layout
 - Creates a figure with appropriate plot which is maximum of 5 plots per row.
- **Boxplot Creation**
 - Generates boxplots for each numeric column in the cleaned dataset
 - Adds red dashed lines for lower bounds and blue dashed lines for upper bounds
 - Includes titles and legend for each subplot
 - Uses `tight_layout()` for proper spacing and displays the plots

```

[41]: # Select numeric columns
numerical_columns = df.select_dtypes(include=['number']).columns

# Create a boolean mask for outliers
outlier_mask = pd.Series(False, index=df.index)

# Calculate IQR and identify outliers
for column in numerical_columns:
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    # Update the outlier mask
    outlier_mask |= (df[column] < lower_bound) | (df[column] > upper_bound)

# Create a new DataFrame without outliers
df_no_outliers = df[~outlier_mask]
df_no_outliers = df_no_outliers.reset_index()
df_no_outliers = df_no_outliers.drop(columns=['index'])

# Set custom color palette
custom_palette = ['#3498db', '#2ecc71', '#e74c3c', '#f1c40f', '#9b59b6',
                  '#1abc9c', '#e67e22', '#34495e', '#7f8c8d', '#16a085']

# Set up the matplotlib figure with a light background
plt.rcParams['figure.facecolor'] = '#f5f5f5'
plt.rcParams['axes.facecolor'] = 'white'

# Create figure
num_cols = 5
num_rows = (len(numerical_columns) + num_cols - 1) // num_cols
fig = plt.figure(figsize=(20, num_rows * 4), dpi=300)

# Create boxplots for the DataFrame without outliers
for i, column in enumerate(numerical_columns):
    Q1 = df_no_outliers[column].quantile(0.25)
    Q3 = df_no_outliers[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    ax = plt.subplot(num_rows, num_cols, i + 1)

    # Create boxplot with custom style
    bp = sns.boxplot(y=df_no_outliers[column],

```

```

        color=custom_palette[i % len(custom_palette)],
        width=0.5,
        linewidth=2,
        fliersize=5,
        showfliers=True)

# Customize the box properties
for box in bp.artists:
    box.set_alpha(0.7)
    box.set_edgecolor('black')

# Add lines for bounds with enhanced style
plt.axhline(y=lower_bound, color='#e74c3c', linestyle='--',
            linewidth=2, label='Lower Bound', alpha=0.8)
plt.axhline(y=upper_bound, color='#3498db', linestyle='--',
            linewidth=2, label='Upper Bound', alpha=0.8)

# Enhanced title and labels
plt.title(column, pad=20, fontsize=12, fontweight='bold')
plt.xlabel('')
plt.ylabel('Value', fontsize=10)

# Add grid for better readability
plt.grid(True, linestyle='--', alpha=0.3)


# Customize spines
for spine in ax.spines.values():
    spine.set_linewidth(1.5)
    spine.set_color('#2c3e50')

# Add the legend with enhanced style
legend = plt.legend(loc='upper right', bbox_to_anchor=(1.2, 1),
                    frameon=True, fancybox=True, shadow=True)

# Adjust tick parameters
plt.tick_params(axis='both', which='major', labelsize=9)

# Adjust layout and display
plt.tight_layout(pad=3.0)
plt.savefig('../output/visuals/features_boxplot.png', dpi=300,
            bbox_inches='tight')
plt.show()

```

time: 11.9 s (started: 2025-01-16 03:21:28 +00:00)

1.2.5 Print summary of outlier removal

```
[42]: print(f"Original dataset shape: {df.shape}")
      print(f"Shape after removing outliers: {df_no_outliers.shape}")
      print(f"Number of rows removed: {df.shape[0] - df_no_outliers.shape[0]}")
      print(f"Outliers detected: {outlier_mask.sum()}")
```

```
Original dataset shape: (47357, 41)
Shape after removing outliers: (20114, 41)
Number of rows removed: 27243
Outliers detected: 27243
time: 2.9 ms (started: 2025-01-16 03:21:40 +00:00)
```

```
[10]: df_no_outliers
```

```
[10]:
```

	index	country	location_name	latitude	longitude	\
0	0	Afghanistan	Kabul	34.5200	69.1800	
1	5	Antigua and Barbuda	Saint John's	17.1200	-61.8500	
2	6	Argentina	Buenos Aires	-34.5900	-58.6700	
3	9	Austria	Vienna	48.2000	16.3700	
4	10	Azerbaijan	Baku	40.4000	49.8800	
...	
20109	47334	Tanzania	Dodoma	-6.1833	35.7500	
20110	47339	Trinidad and Tobago	Port Of Spain	10.6500	-61.5167	
20111	47349	Uruguay	Montevideo	-34.8581	-56.1708	
20112	47354	Yemen	Sanaa	15.3547	44.2067	
20113	47355	Zambia	Lusaka	-15.4167	28.2833	

	timezone	last_updated_epoch	last_updated	\
0	Asia/Kabul	1715849100	2024-05-16 13:15	
1	America/Antigua	1715849100	2024-05-16 04:45	
2	America/Argentina/Buenos_Aires	1715849100	2024-05-16 05:45	
3	Europe/Vienna	1715849100	2024-05-16 10:45	
4	Asia/Baku	1715849100	2024-05-16 12:45	
...	
20109	Africa/Dar_es_Salaam	1736937000	2025-01-15 13:30	
20110	America/Port_of_Spain	1736937000	2025-01-15 06:30	
20111	America/Montevideo	1736940600	2025-01-15 08:30	
20112	Asia/Aden	1736937000	2025-01-15 13:30	
20113	Africa/Lusaka	1736937000	2025-01-15 12:30	

	temperature_celsius	temperature_fahrenheit	...	air_quality_PM2.5	\
0	26.6	79.8	...	8.400	
1	26.0	78.8	...	1.200	
2	8.0	46.4	...	4.000	
3	16.0	60.8	...	3.700	
4	17.0	62.6	...	1.900	
...	

20109	27.2	81.0	...	13.320
20110	22.4	72.3	...	10.915
20111	22.2	72.0	...	10.545
20112	18.9	66.0	...	26.455
20113	28.0	82.4	...	15.355

	air_quality_PM10	air_quality_us-epa-index	air_quality_gb-defra-index	\
0	26.600	1		1
1	4.500	1		1
2	5.300	1		1
3	4.400	1		1
4	2.200	1		1
...	
20109	15.540	1		2
20110	15.910	1		1
20111	13.135	1		1
20112	69.745	2		3
20113	15.355	1		2

	sunrise	sunset	moonrise	moonset	moon_phase	\
0	04:50 AM	06:50 PM	12:12 PM	01:11 AM	Waxing Gibbous	
1	05:36 AM	06:32 PM	01:05 PM	01:14 AM	Waxing Gibbous	
2	07:43 AM	05:59 PM	02:36 PM	01:04 AM	Waxing Gibbous	
3	05:14 AM	08:29 PM	01:00 PM	02:42 AM	Waxing Gibbous	
4	05:23 AM	07:51 PM	12:54 PM	02:10 AM	Waxing Gibbous	
...	
20109	06:33 AM	07:00 PM	08:29 PM	07:45 AM	Waning Gibbous	
20110	06:28 AM	06:03 PM	07:48 PM	07:53 AM	Waning Gibbous	
20111	05:48 AM	08:01 PM	09:32 PM	07:14 AM	Waning Gibbous	
20112	06:33 AM	05:52 PM	07:23 PM	07:45 AM	Waning Gibbous	
20113	05:48 AM	06:44 PM	08:13 PM	07:00 AM	Waning Gibbous	

	moon_illumination
0	55
1	55
2	55
3	55
4	55
...	...
20109	99
20110	98
20111	98
20112	99
20113	99

[20114 rows x 42 columns]

time: 23.8 ms (started: 2025-01-16 02:52:34 +00:00)

1.3 Exploratory Data Analysis (EDA)

1.3.1 Get summary statistics

```
[11]: print(df_no_outliers.describe())
```

	index	latitude	longitude	last_updated_epoch \
count	20114.000000	20114.000000	20114.000000	2.011400e+04
mean	21088.637317	18.757638	9.237831	1.725244e+09
std	12880.075996	25.750078	47.464237	5.788978e+06
min	0.000000	-34.860000	-90.530000	1.715849e+09
25%	10352.750000	-0.216700	-9.130000	1.720357e+09
50%	19926.500000	15.354700	18.050000	1.724762e+09
75%	30929.750000	42.000000	33.780000	1.729675e+09
max	47355.000000	63.830000	134.557800	1.736941e+09

	temperature_celsius	temperature_fahrenheit	wind_mph \
count	20114.000000	20114.000000	20114.000000
mean	25.080392	77.145913	8.491136
std	7.080302	12.744840	4.665866
min	2.500000	36.500000	2.200000
25%	21.300000	70.300000	4.500000
50%	26.300000	79.300000	8.100000
75%	29.400000	84.900000	11.900000
max	44.400000	112.000000	22.800000

	wind_kph	wind_degree	pressure_mb ...	gust_kph \
count	20114.000000	20114.000000	20114.000000 ...	20114.000000
mean	13.670026	168.490554	1013.865218 ...	18.811738
std	7.507223	103.028148	5.123220 ...	9.382779
min	3.600000	1.000000	1000.000000 ...	3.600000
25%	7.200000	81.000000	1011.000000 ...	10.900000
50%	13.000000	160.000000	1013.000000 ...	17.400000
75%	19.100000	255.000000	1017.000000 ...	25.100000
max	36.700000	360.000000	1027.000000 ...	48.200000

	air_quality_Carbon_Monoxide	air_quality_Ozone \
count	20114.000000	20114.000000
mean	304.738194	69.270707
std	134.286664	31.155491
min	93.000000	0.000000
25%	205.350000	47.000000
50%	262.700000	67.200000
75%	360.750000	91.600000
max	847.800000	153.100000

	air_quality_Nitrogen_dioxide	air_quality_Sulphur_dioxide	\
count	20114.000000	20114.000000	
mean	4.265833	2.798284	
std	6.082915	3.457773	
min	0.000000	0.000000	
25%	0.696250	0.555000	
50%	1.700000	1.431000	
75%	4.810000	3.700000	
max	32.930000	17.945000	

	air_quality_PM2.5	air_quality_PM10	air_quality_us-epa-index	\
count	20114.000000	20114.000000	20114.000000	
mean	10.890141	18.688152	1.261062	
std	9.434208	17.773722	0.469111	
min	0.500000	0.500000	1.000000	
25%	3.694000	6.000000	1.000000	
50%	8.300000	13.300000	1.000000	
75%	15.355000	24.900000	1.000000	
max	53.896000	92.315000	3.000000	

	air_quality_gb-defra-index	moon_illumination
count	20114.000000	20114.000000
mean	1.490156	48.641792
std	0.803666	35.011318
min	1.000000	0.000000
25%	1.000000	13.000000
50%	1.000000	49.000000
75%	2.000000	83.000000
max	6.000000	100.000000

[8 rows x 31 columns]

time: 50.9 ms (started: 2025-01-16 02:52:34 +00:00)

[12]: df_no_outliers.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 20114 entries, 0 to 20113

Data columns (total 42 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	index	20114 non-null	int64
1	country	20114 non-null	object
2	location_name	20114 non-null	object
3	latitude	20114 non-null	float64
4	longitude	20114 non-null	float64
5	timezone	20114 non-null	object
6	last_updated_epoch	20114 non-null	int64
7	last_updated	20114 non-null	object

8	temperature_celsius	20114	non-null	float64
9	temperature_fahrenheit	20114	non-null	float64
10	condition_text	20114	non-null	object
11	wind_mph	20114	non-null	float64
12	wind_kph	20114	non-null	float64
13	wind_degree	20114	non-null	int64
14	wind_direction	20114	non-null	object
15	pressure_mb	20114	non-null	float64
16	pressure_in	20114	non-null	float64
17	precip_mm	20114	non-null	float64
18	precip_in	20114	non-null	float64
19	humidity	20114	non-null	int64
20	cloud	20114	non-null	int64
21	feels_like_celsius	20114	non-null	float64
22	feels_like_fahrenheit	20114	non-null	float64
23	visibility_km	20114	non-null	float64
24	visibility_miles	20114	non-null	float64
25	uv_index	20114	non-null	float64
26	gust_mph	20114	non-null	float64
27	gust_kph	20114	non-null	float64
28	air_quality_Carbon_Monoxide	20114	non-null	float64
29	air_quality_Ozone	20114	non-null	float64
30	air_quality_Nitrogen_dioxide	20114	non-null	float64
31	air_quality_Sulphur_dioxide	20114	non-null	float64
32	air_quality_PM2.5	20114	non-null	float64
33	air_quality_PM10	20114	non-null	float64
34	air_quality_us-epa-index	20114	non-null	int64
35	air_quality_gb-defra-index	20114	non-null	int64
36	sunrise	20114	non-null	object
37	sunset	20114	non-null	object
38	moonrise	20114	non-null	object
39	moonset	20114	non-null	object
40	moon_phase	20114	non-null	object
41	moon_illumination	20114	non-null	int64

dtypes: float64(23), int64(8), object(11)

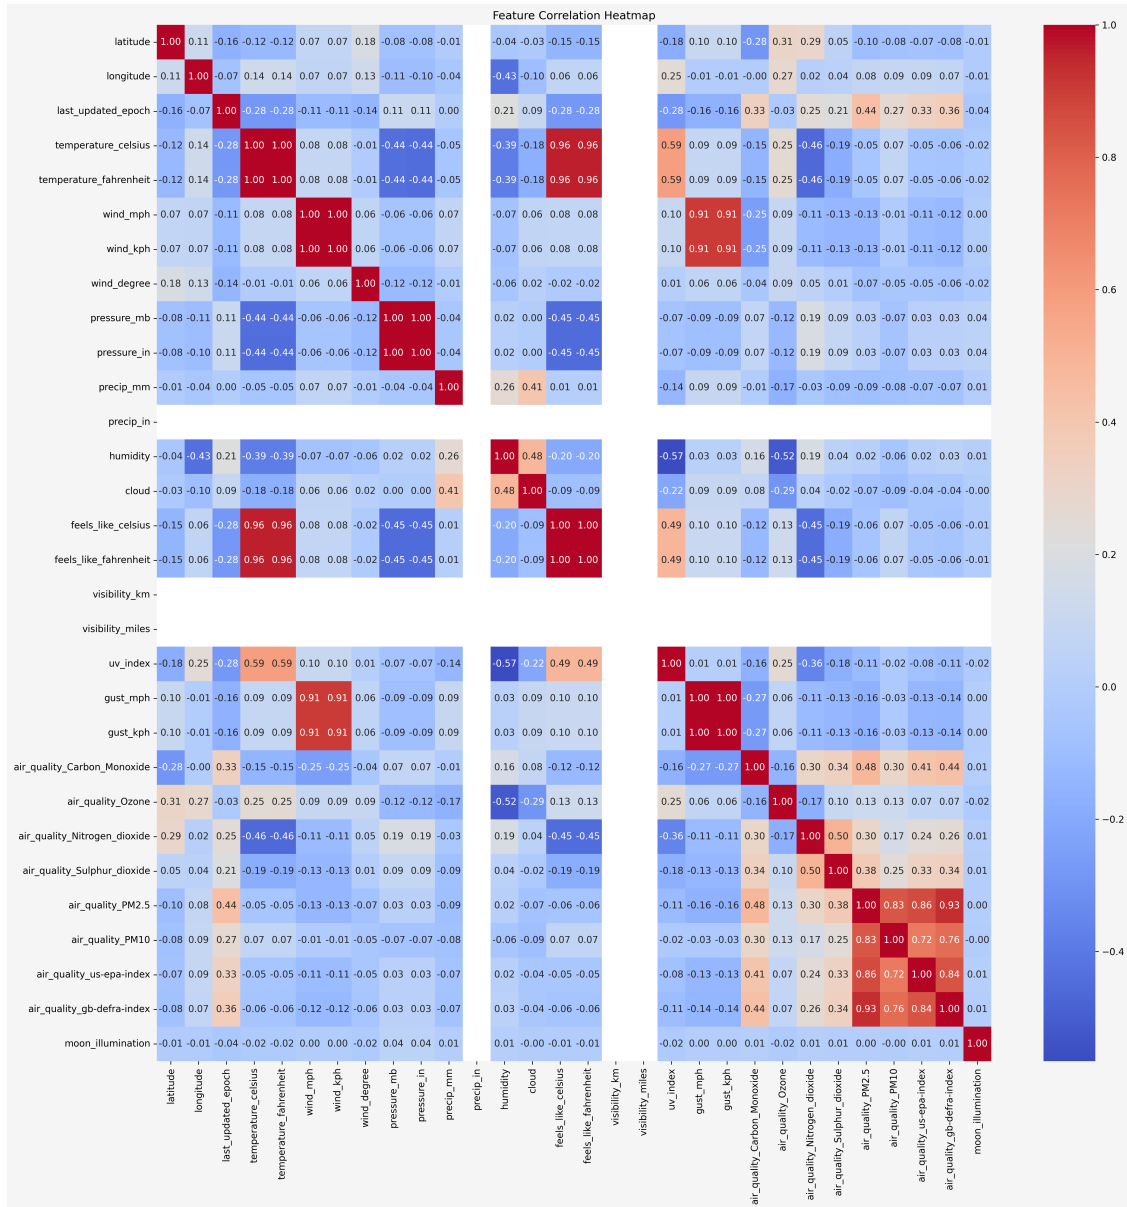
memory usage: 6.4+ MB

time: 10.9 ms (started: 2025-01-16 02:52:34 +00:00)

1.3.2 Correlation Heatmap

```
[13]: df_numeric = df_no_outliers[numerical_columns]
plt.figure(figsize=(20, 20), dpi=300)
sns.heatmap(df_numeric.corr(), annot=True, cmap="coolwarm", fmt=".2f",
            cbar=True, linewidths=0, linecolor='white')
plt.title("Feature Correlation Heatmap")
plt.grid(False)
plt.savefig('../output/visuals/corr_heatmap.png', dpi=300, bbox_inches='tight')
```

```
plt.show()
```



time: 4.55 s (started: 2025-01-16 02:52:34 +00:00)

1.3.3 Use 'last_updated' as index

use last_update as index time series and drop unnecessary columns such as 'index', 'last_updated_epoch'

```
[14]: # Convert 'last_updated' to datetime format
df_no_outliers['last_updated'] = pd.to_datetime(df_no_outliers['last_updated'])
```

```

# Extracting temporal features
df_no_outliers['year'] = df_no_outliers['last_updated'].dt.year
df_no_outliers['month'] = df_no_outliers['last_updated'].dt.month
df_no_outliers['day'] = df_no_outliers['last_updated'].dt.day
df_no_outliers['hour'] = df_no_outliers['last_updated'].dt.hour

# Set 'last_updated' as the index
df_no_outliers.set_index('last_updated', inplace=True)

# Sort by index if needed
df_no_outliers.sort_index(inplace=True)

# Drop columns
df_no_outliers = df_no_outliers.drop(columns=['last_updated_epoch'])

# Display the result
df_no_outliers

```

```

[14]:
      last_updated      country  location_name  latitude  longitude \
2024-05-16 02:45:00  Nicaragua      Managua    12.1500   -86.2700
2024-05-16 02:45:00    Belize      Belmopan   17.2500   -88.7700
2024-05-16 03:45:00    Panama    Panama City    8.9700   -79.5300
2024-05-16 03:45:00     Peru        Lima   -12.0500   -77.0500
2024-05-16 04:45:00   Grenada  Saint George's   12.0500   -61.7500
...
2025-01-15 15:15:00  Maldives    Dhidhdhoo    6.8833    73.1000
2025-01-15 15:45:00    Bhutan     Thimphu    27.4833    89.6000
2025-01-15 16:45:00  Thailand         Nan    18.7833   100.7833
2025-01-15 17:30:00  Cambodia    Phnom Penh   11.5500   104.9167
2025-01-15 19:15:00  Philippines      Manila   14.6042   120.9822

      last_updated      timezone  temperature_celsius \
2024-05-16 02:45:00  America/Managua          27.2
2024-05-16 02:45:00  America/Belize           26.0
2024-05-16 03:45:00  America/Panama           26.0
2024-05-16 03:45:00  America/Lima             16.6
2024-05-16 04:45:00  America/Grenada          28.0
...
2025-01-15 15:15:00  Indian/Maldives          27.2
2025-01-15 15:45:00  Asia/Thimphu             11.0
2025-01-15 16:45:00  Asia/Bangkok             28.2
2025-01-15 17:30:00  Asia/Phnom_Penh          31.2
2025-01-15 19:15:00  Asia/Manila              26.6

```


	temperature_fahrenheit	condition_text	wind_mph	\
last_updated				
2024-05-16 02:45:00	80.9	Patchy rain nearby	3.6	
2024-05-16 02:45:00	78.9	Overcast	4.3	
2024-05-16 03:45:00	78.8	Overcast	2.2	
2024-05-16 03:45:00	61.9	Partly Cloudy	7.4	
2024-05-16 04:45:00	82.4	Partly cloudy	13.6	
...	
2025-01-15 15:15:00	80.9	Overcast	14.5	
2025-01-15 15:45:00	51.7	Sunny	5.8	
2025-01-15 16:45:00	82.8	Sunny	2.5	
2025-01-15 17:30:00	88.2	Partly cloudy	8.9	
2025-01-15 19:15:00	79.9	Patchy rain nearby	7.2	

	wind_kph	...	sunrise	sunset	moonrise	moonset	\
last_updated		...					
2024-05-16 02:45:00	5.8	...	05:21 AM	06:02 PM	12:49 PM	12:49 AM	
2024-05-16 02:45:00	6.8	...	05:23 AM	06:20 PM	12:56 PM	01:04 AM	
2024-05-16 03:45:00	3.6	...	05:58 AM	06:31 PM	01:24 PM	01:18 AM	
2024-05-16 03:45:00	11.9	...	06:18 AM	05:51 PM	01:30 PM	12:47 AM	
2024-05-16 04:45:00	22.0	...	05:43 AM	06:24 PM	01:08 PM	01:08 AM	
...	
2025-01-15 15:15:00	23.4	...	06:24 AM	06:10 PM	07:36 PM	07:31 AM	
2025-01-15 15:45:00	9.4	...	06:53 AM	05:29 PM	06:53 PM	08:01 AM	
2025-01-15 16:45:00	4.0	...	06:53 AM	06:00 PM	07:23 PM	07:58 AM	
2025-01-15 17:30:00	14.4	...	06:24 AM	05:56 PM	07:17 PM	07:28 AM	
2025-01-15 19:15:00	11.5	...	06:25 AM	05:46 PM	07:05 PM	07:27 AM	

	moon_phase	moon_illumination	year	month	day	hour
last_updated						
2024-05-16 02:45:00	Waxing Gibbous	55	2024	5	16	2
2024-05-16 02:45:00	Waxing Gibbous	55	2024	5	16	2
2024-05-16 03:45:00	Waxing Gibbous	55	2024	5	16	3
2024-05-16 03:45:00	Waxing Gibbous	55	2024	5	16	3
2024-05-16 04:45:00	Waxing Gibbous	55	2024	5	16	4
...
2025-01-15 15:15:00	Waning Gibbous	99	2025	1	15	15
2025-01-15 15:45:00	Waning Gibbous	99	2025	1	15	15
2025-01-15 16:45:00	Waning Gibbous	99	2025	1	15	16
2025-01-15 17:30:00	Waning Gibbous	99	2025	1	15	17
2025-01-15 19:15:00	Waning Gibbous	99	2025	1	15	19

[20114 rows x 43 columns]

time: 38.6 ms (started: 2025-01-16 02:52:38 +00:00)

1.3.4 Daily Global Temperature (Celsius) Trends

```
[15]: # Filter for object columns
object_columns = df_no_outliers.select_dtypes(include='object')

# Resample to daily frequency by taking the mean of the temperature for each day
df_daily = df_no_outliers.drop(columns=object_columns)

df_daily = df_daily.resample('D').mean()

# Calculate the mean of the temperature celcius series (ignoring NaN)
mean_value = df_daily['temperature_celsius'].mean()

# Replace the specific value 7.1 with the mean
df_daily = df_daily.replace(7.1, mean_value)

# Forward fill
df_daily = df_daily.fillna(method='ffill')

# Calculate the rolling average
rolling_window = 30 # Set the window size (e.g., 30 days for a monthly average)
df_daily['rolling_avg'] = df_daily['temperature_celsius'].
    ↪rolling(window=rolling_window).mean()

# Plot the overall temperature trend with the rolling average
plt.figure(figsize=(12, 5))
sns.set(style="whitegrid")

# Plot the overall trend
plt.plot(
    df_daily.index,
    df_daily['temperature_celsius'],
    label="Overall Trend",
    alpha=0.6,
    color="gray"
)

# Plot the rolling average
plt.plot(
    df_daily.index,
    df_daily['rolling_avg'],
    label=f"{rolling_window}-Day Moving Average",
    color="blue",
    linewidth=2
)


# Add plot details
```

```
plt.title("Overall Global Temperature Trend Over Time", fontsize=16)
plt.xlabel("Time", fontsize=12)
plt.ylabel("Temperature (°C)", fontsize=12)
plt.xticks(rotation=45)
plt.legend() # Add legend for clarity
plt.tight_layout()

# Show the plot
plt.savefig('../output/visuals/daily_global_temperature_trends.png', dpi=300,
            bbox_inches='tight')
plt.show()
```

/tmp/ipykernel_5464/2681601598.py:16: FutureWarning: DataFrame.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.

```
df_daily = df_daily.fillna(method='ffill')
```



time: 740 ms (started: 2025-01-16 02:52:38 +00:00)

1.3.5 HeatMap Location based on Temperature

```
[16]: # Create a Folium map centered at a specific latitude and longitude
m = folium.Map(location=[0, 0], zoom_start=2) # You can adjust the coordinates
      and zoom level

# Create a list of coordinates and corresponding values (e.g., temperature)
locations = df_no_outliers[['latitude', 'longitude']].values
values = df_no_outliers['temperature_celsius'].values

# Normalize the values for the heatmap (adjust as needed)
max_value = max(values)
```

```

normalized_values = [v / max_value for v in values]

# Create a HeatMap layer on the map
HeatMap(list(zip(locations[:, 0], locations[:, 1], normalized_values))).
    ↪add_to(m)

# Display the map in the jupyter notebook
display(m)

```

<folium.folium.Map at 0x7f86da5dd550>

time: 370 ms (started: 2025-01-16 02:52:39 +00:00)

1.3.6 Daily Global Precipitation Trends

```

[17]: # Filter for object columns
object_columns = df_no_outliers.select_dtypes(include='object')

# Resample to daily frequency by taking the mean of the temperature for each day
df_daily = df_no_outliers.drop(columns=object_columns)

df_daily = df_daily.resample('D').mean()

# Calculate the mean of the temperature celcius series (ignoring NaN)
mean_value = df_daily['precip_mm'].mean()

# Replace the specific value 7.1 with the mean
df_daily = df_daily.replace(7.1, mean_value)

# Forward fill
df_daily = df_daily.fillna(method='ffill')

# Calculate the rolling average
rolling_window = 30 # Set the window size (e.g., 30 days for a monthly average)
rolling_avg = df_daily['precip_mm'].rolling(window=rolling_window).mean()

# Plot the overall temperature trend with the rolling average
plt.figure(figsize=(12, 5))
sns.set(style="whitegrid")

# Plot the overall trend
plt.plot(
    df_daily.index,
    df_daily['precip_mm'],
    label="Overall Trend",
    alpha=0.6,
    color="grey"
)

```

```

# Plot the rolling average
plt.plot(
    df_daily.index,
    rolling_avg,
    label=f"{rolling_window}-Day Moving Average",
    color="orange",
    linewidth=2
)


# Add plot details
plt.title("Overall Global Precipitation Trend Over Time", fontsize=16)
plt.xlabel("Time", fontsize=12)
plt.ylabel("Precipitation", fontsize=12)
plt.xticks(rotation=45)
plt.legend() # Add legend for clarity
plt.tight_layout()

# Show the plot
plt.savefig('../output/visuals/daily_global_precipitation_trends.png', dpi=300,
            bbox_inches='tight')
plt.show()

```

/tmp/ipykernel_5464/3249255930.py:16: FutureWarning: DataFrame.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.

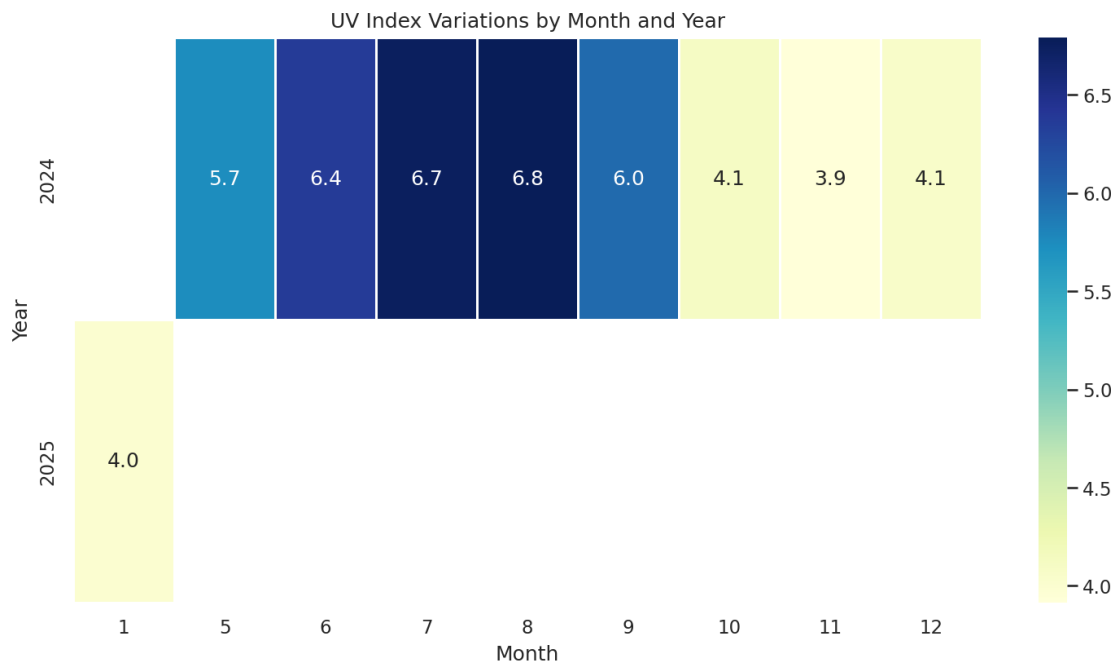
```
df_daily = df_daily.fillna(method='ffill')
```



time: 775 ms (started: 2025-01-16 02:52:40 +00:00)

1.3.7 UV Index Analysis

```
[18]: # Heatmap to visualize UV index variations by month and year
uv_index_heatmap = df_no_outliers.pivot_table(values='uv_index', index='year',
columns='month', aggfunc='mean')
plt.figure(figsize=(12, 6), dpi=150)
sns.heatmap(uv_index_heatmap, cmap='YlGnBu', annot=True, fmt='.1f',
linewidths=0.5)
plt.title('UV Index Variations by Month and Year')
plt.xlabel('Month')
plt.ylabel('Year')
plt.grid(False)
plt.savefig('../output/visuals/UV_index.png', dpi=300, bbox_inches='tight')
plt.show()
```



time: 537 ms (started: 2025-01-16 02:52:40 +00:00)

1.4 Feature Engineering

Search for all 'object' data and convert it into numerical data so that it can be used as a feature in machine learning models.

```
[19]: # Filter for object columns
object_columns = df_no_outliers.select_dtypes(include='object')

# Print unique values for each object column
```

```
for col in object_columns.columns:
    print(f"Unique values in column '{col}':")
```

```
Unique values in column 'country':
Unique values in column 'location_name':
Unique values in column 'timezone':
Unique values in column 'condition_text':
Unique values in column 'wind_direction':
Unique values in column 'sunrise':
Unique values in column 'sunset':
Unique values in column 'moonrise':
Unique values in column 'moonset':
Unique values in column 'moon_phase':
time: 5.19 ms (started: 2025-01-16 02:52:41 +00:00)
```

Preprocess country, location_name, condition, moon phase, and wind direction by converting text to lowercase, then encoding unique values into numeric codes (0,1,2,...)

```
[20]: # Convert to lowercase
df_no_outliers['condition_text'] = df_no_outliers['condition_text'].str.lower()
df_no_outliers['moon_phase'] = df_no_outliers['moon_phase'].str.lower()
df_no_outliers['wind_direction'] = df_no_outliers['wind_direction'].str.lower()
df_no_outliers['country'] = df_no_outliers['country'].str.lower()
df_no_outliers['location_name'] = df_no_outliers['location_name'].str.lower()
df_no_outliers['timezone'] = df_no_outliers['timezone'].str.lower()

# Get unique values and assign numeric codes
df_no_outliers['condition_text'] = df_no_outliers['condition_text'].
    ↪astype('category').cat.codes
df_no_outliers['moon_phase'] = df_no_outliers['moon_phase'].astype('category').
    ↪cat.codes
df_no_outliers['wind_direction'] = df_no_outliers['wind_direction'].
    ↪astype('category').cat.codes
df_no_outliers['country'] = df_no_outliers['country'].astype('category').cat.
    ↪codes
df_no_outliers['location_name'] = df_no_outliers['location_name'].
    ↪astype('category').cat.codes
df_no_outliers['timezone'] = df_no_outliers['timezone'].astype('category').cat.
    ↪codes
```

```
time: 26.6 ms (started: 2025-01-16 02:52:41 +00:00)
```

Preprocess all time features into numerical data

```
[21]: def time_to_minutes(time_str):
        if time_str.lower() == 'no moonrise' or time_str.lower() == 'no moonset' or
        ↪time_str.lower() == 'no sunrise' or time_str.lower() == 'no sunset':
```

```

        return -1 # or some other placeholder value (e.g., -1)
    try:
        time_obj = pd.to_datetime(time_str, format='%I:%M %p') # Convert to
        ↪datetime
        return time_obj.hour * 60 + time_obj.minute
    except ValueError:
        return None # or handle other invalid time formats similarly

df_no_outliers['moonrise'] = df_no_outliers['moonrise'].apply(time_to_minutes)
df_no_outliers['moonset'] = df_no_outliers['moonset'].apply(time_to_minutes)
df_no_outliers['sunrise'] = df_no_outliers['sunrise'].apply(time_to_minutes)
df_no_outliers['sunset'] = df_no_outliers['sunset'].apply(time_to_minutes)

```

time: 5.37 s (started: 2025-01-16 02:52:41 +00:00)

[22]: df_no_outliers

[22]:

	country	location_name	latitude	longitude	timezone	\
last_updated						
2024-05-16 02:45:00	111	110	12.1500	-86.2700	62	
2024-05-16 02:45:00	16	36	17.2500	-88.7700	47	
2024-05-16 03:45:00	119	141	8.9700	-79.5300	67	
2024-05-16 03:45:00	121	99	-12.0500	-77.0500	61	
2024-05-16 04:45:00	63	164	12.0500	-61.7500	54	
...	
2025-01-15 15:15:00	96	64	6.8833	73.1000	163	
2025-01-15 15:45:00	18	185	27.4833	89.6000	113	
2025-01-15 16:45:00	155	129	18.7833	100.7833	84	
2025-01-15 17:30:00	27	146	11.5500	104.9167	103	
2025-01-15 19:15:00	122	112	14.6042	120.9822	101	

	temperature_celsius	temperature_fahrenheit	\
last_updated			
2024-05-16 02:45:00	27.2	80.9	
2024-05-16 02:45:00	26.0	78.9	
2024-05-16 03:45:00	26.0	78.8	
2024-05-16 03:45:00	16.6	61.9	
2024-05-16 04:45:00	28.0	82.4	
...	
2025-01-15 15:15:00	27.2	80.9	
2025-01-15 15:45:00	11.0	51.7	
2025-01-15 16:45:00	28.2	82.8	
2025-01-15 17:30:00	31.2	88.2	
2025-01-15 19:15:00	26.6	79.9	

	condition_text	wind_mph	wind_kph	...	sunrise	sunset	\
last_updated				...			

2024-05-16 02:45:00	14	3.6	5.8	...	321	1082
2024-05-16 02:45:00	10	4.3	6.8	...	323	1100
2024-05-16 03:45:00	10	2.2	3.6	...	358	1111
2024-05-16 03:45:00	11	7.4	11.9	...	378	1071
2024-05-16 04:45:00	11	13.6	22.0	...	343	1104
...
2025-01-15 15:15:00	10	14.5	23.4	...	384	1090
2025-01-15 15:45:00	16	5.8	9.4	...	413	1049
2025-01-15 16:45:00	16	2.5	4.0	...	413	1080
2025-01-15 17:30:00	11	8.9	14.4	...	384	1076
2025-01-15 19:15:00	14	7.2	11.5	...	385	1066

	moonrise	moonset	moon_phase	moon_illumination	year	\
last_updated						
2024-05-16 02:45:00	769	49	7		55	2024
2024-05-16 02:45:00	776	64	7		55	2024
2024-05-16 03:45:00	804	78	7		55	2024
2024-05-16 03:45:00	810	47	7		55	2024
2024-05-16 04:45:00	788	68	7		55	2024
...
2025-01-15 15:15:00	1176	451	5		99	2025
2025-01-15 15:45:00	1133	481	5		99	2025
2025-01-15 16:45:00	1163	478	5		99	2025
2025-01-15 17:30:00	1157	448	5		99	2025
2025-01-15 19:15:00	1145	447	5		99	2025

	month	day	hour
last_updated			
2024-05-16 02:45:00	5	16	2
2024-05-16 02:45:00	5	16	2
2024-05-16 03:45:00	5	16	3
2024-05-16 03:45:00	5	16	3
2024-05-16 04:45:00	5	16	4
...
2025-01-15 15:15:00	1	15	15
2025-01-15 15:45:00	1	15	15
2025-01-15 16:45:00	1	15	16
2025-01-15 17:30:00	1	15	17
2025-01-15 19:15:00	1	15	19

[20114 rows x 43 columns]

time: 15.4 ms (started: 2025-01-16 02:52:46 +00:00)

1.5 Data Normalization

```
[23]: # Move temperature_celciuse into last

cols = [col for col in df_no_outliers.columns if col != 'temperature_celsius']
    ↪ # All columns except 'temperature_celsius'
df_no_outliers = df_no_outliers[cols + ['temperature_celsius']] # Reorder
    ↪ columns, placing 'temperature_celsius' at the end
```

time: 2.87 ms (started: 2025-01-16 02:52:46 +00:00)

```
[24]: # Define the scale pipeline
features = df_no_outliers.columns
scaler = ColumnTransformer(
    transformers=[('scaler', RobustScaler(), features)]
)

# Fit and transform the data
values = scaler.fit_transform(df_no_outliers)
```

time: 34.6 ms (started: 2025-01-16 02:52:46 +00:00)

1.6 Data Preprocessing

```
[25]: n_steps_in = 120
n_steps_out = 1

def preprocess_data(values_array, n_steps_in=14, n_steps_out=5, train_split=0.
    ↪ 8):
    """
    Preprocess a single dataset for training.

    Args:
        values_array: Numpy array containing time series data with features
        n_steps_in: Number of lookback days
        n_steps_out: Number of prediction days
        train_split: Train/validation split ratio

    Returns:
        tuple: (train_X, train_y, val_X, val_y, global_scaler)
    """

    # 1. Prepare data
    # Remove 'Close Next Day' from features (last column)
    features = values_array[:, :-1] # All columns except the last one
    targets = values_array[:, -1]   # Only the last column

    # 2. Split into train/validation
```

```

n_train = int(len(features) * train_split)

# Ensure we have enough data for both training and validation
if n_train <= n_steps_in + n_steps_out:
    raise ValueError(f"Insufficient data. Need more than {n_steps_in + \
↪n_steps_out} samples.")

# Split features and targets
train_features = features[:n_train]
train_targets = targets[:n_train]
val_features = features[n_train:]
val_targets = targets[n_train:]

# 4. Create sequences
train_X, train_y = create_sequences(train_features, train_targets, \
↪n_steps_in, n_steps_out)
val_X, val_y = create_sequences(val_features, val_targets, n_steps_in, \
↪n_steps_out)

print(f"Training shapes: X={train_X.shape}, y={train_y.shape}")
print(f"Validation shapes: X={val_X.shape}, y={val_y.shape}")
print(f"Number of features: {train_X.shape[2]}")

return train_X, train_y, val_X, val_y

def create_sequences(features, targets, n_steps_in, n_steps_out):
    """
    Generate synchronized sequences for LSTM input features and output targets.

    Args:
        features: Scaled feature data (numpy array)
        targets: Scaled target data (numpy array)
        n_steps_in: Number of input time steps
        n_steps_out: Number of output time steps

    Returns:
        tuple: (X sequences, y sequences)
    """
    X, y = [], []

    # Ensure we have enough data for sequence creation
    if len(features) < n_steps_in + n_steps_out:
        raise ValueError("Data length is too short for the specified sequence \
↪lengths")

    for i in range(len(features) - n_steps_in - n_steps_out + 1):
        # Input sequence (n_steps_in days of all features)

```

```

seq_x = features[i:(i + n_steps_in)]
# Output sequence (next n_steps_out days of target variable)
seq_y = targets[(i + n_steps_in):(i + n_steps_in + n_steps_out)]

X.append(seq_x)
y.append(seq_y)

return np.array(X), np.array(y)

train_X, train_y, val_X, val_y = preprocess_data(values, n_steps_in,
↪n_steps_out, train_split=0.8)

```

Training shapes: X=(15971, 120, 42), y=(15971, 1)
 Validation shapes: X=(3903, 120, 42), y=(3903, 1)
 Number of features: 42
 time: 3.45 s (started: 2025-01-16 02:52:46 +00:00)

1.7 Model Training

Models Overview

1. LSTM Model:

- Sequence-based deep learning model.
- Architecture includes an LSTM layer (32 units) and a Dense output layer.
- Optimized using Adam optimizer and MSE as the loss function.
- Early stopping applied to prevent overfitting.

2. GRU Model:

- Similar structure to the LSTM model but uses a GRU layer instead.
- Early stopping used for efficient training.

3. XGBoost Model:

- Gradient boosting algorithm for regression tasks.
- Input data reshaped into 2D format.
- Key hyperparameters: 100 estimators, max depth of 6, learning rate 0.1, and subsample ratios of 0.8.

```

[26]: tf.keras.backend.clear_session()

def build_lstm_model(n_steps_in, n_features, n_steps_out):
    model = Sequential([
        LSTM(32, activation='relu', input_shape=(n_steps_in, n_features)),
        Dense(n_steps_out)
    ])

    optimizer = tf.keras.optimizers.Adam(1e-4)
    model.compile(optimizer=optimizer, loss=['mse'])
    model.summary()

    return model

```

```
# Create and train model
```

```
lstm_model = build_lstm_model(n_steps_in=n_steps_in, n_features=train_X.  
    ↪shape[2], n_steps_out=n_steps_out)
```

```
2025-01-16 02:52:50.445556: I  
external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful  
NUMA node read from SysFS had negative value (-1), but there must be at least  
one NUMA node, so returning NUMA node zero. See more at  
https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-  
pci#L344-L355
```

```
2025-01-16 02:52:50.762310: I  
external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful  
NUMA node read from SysFS had negative value (-1), but there must be at least  
one NUMA node, so returning NUMA node zero. See more at  
https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-  
pci#L344-L355
```

```
2025-01-16 02:52:50.762507: I  
external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful  
NUMA node read from SysFS had negative value (-1), but there must be at least  
one NUMA node, so returning NUMA node zero. See more at  
https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-  
pci#L344-L355
```

WARNING:tensorflow:Layer lstm will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.

```
2025-01-16 02:52:50.764435: I  
external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful  
NUMA node read from SysFS had negative value (-1), but there must be at least  
one NUMA node, so returning NUMA node zero. See more at  
https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-  
pci#L344-L355
```

```
2025-01-16 02:52:50.764628: I  
external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful  
NUMA node read from SysFS had negative value (-1), but there must be at least  
one NUMA node, so returning NUMA node zero. See more at  
https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-  
pci#L344-L355
```

```
2025-01-16 02:52:50.764717: I  
external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful  
NUMA node read from SysFS had negative value (-1), but there must be at least  
one NUMA node, so returning NUMA node zero. See more at  
https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-  
pci#L344-L355
```

```
2025-01-16 02:52:52.928331: I  
external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful  
NUMA node read from SysFS had negative value (-1), but there must be at least
```

one NUMA node, so returning NUMA node zero. See more at
<https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355>

2025-01-16 02:52:52.928509: I

external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful
NUMA node read from SysFS had negative value (-1), but there must be at least
one NUMA node, so returning NUMA node zero. See more at

<https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355>

2025-01-16 02:52:52.928618: I

external/local_xla/xla/stream_executor/cuda/cuda_executor.cc:901] successful
NUMA node read from SysFS had negative value (-1), but there must be at least
one NUMA node, so returning NUMA node zero. See more at

<https://github.com/torvalds/linux/blob/v6.0/Documentation/ABI/testing/sysfs-bus-pci#L344-L355>

2025-01-16 02:52:52.928696: I

tensorflow/core/common_runtime/gpu/gpu_device.cc:1929] Created device
/job:localhost/replica:0/task:0/device:GPU:0 with 14223 MB memory: -> device:
0, name: NVIDIA RTX A4000, pci bus id: 0000:00:05.0, compute capability: 8.6

Model: "sequential"

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 32)	9600
dense (Dense)	(None, 1)	33

Total params: 9633 (37.63 KB)

Trainable params: 9633 (37.63 KB)

Non-trainable params: 0 (0.00 Byte)

time: 3.3 s (started: 2025-01-16 02:52:50 +00:00)

```
[27]: early_stopping = tf.keras.callbacks.EarlyStopping(  
    monitor='val_loss',  
    patience=10,  
    mode='min',  
    restore_best_weights=True  
)  
  
history = lstm_model.fit(  
    train_X,  
    train_y,  
    validation_data=(val_X, val_y),  
    epochs=100,  
    batch_size=32,
```

```

        verbose=2,
        shuffle=False,
        callbacks=[early_stopping],
    )

    lstm_model.save('../output/model/lstm_model.keras')

```

Epoch 1/100

2025-01-16 02:53:07.059219: I external/local_xla/xla/service/service.cc:168] XLA service 0x7f85ac0131e0 initialized for platform CUDA (this does not guarantee that XLA will be used). Devices:

2025-01-16 02:53:07.059318: I external/local_xla/xla/service/service.cc:176] StreamExecutor device (0): NVIDIA RTX A4000, Compute Capability 8.6

2025-01-16 02:53:07.087717: I tensorflow/compiler/mlir/tensorflow/utils/dump_mlir_util.cc:269] disabling MLIR crash reproducer, set env var `MLIR_CRASH_REPRODUCER_DIRECTORY` to enable.

2025-01-16 02:53:08.022140: I external/local_xla/xla/stream_executor/cuda/cuda_dnn.cc:454] Loaded cuDNN version 8907

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1736995988.167958 5524 device_compiler.h:186] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.

500/500 - 47s - loss: 0.5678 - val_loss: 1.1067 - 47s/epoch - 95ms/step

Epoch 2/100

500/500 - 38s - loss: 0.5551 - val_loss: 1.0865 - 38s/epoch - 75ms/step

Epoch 3/100

500/500 - 38s - loss: 0.5489 - val_loss: 1.0825 - 38s/epoch - 76ms/step

Epoch 4/100

500/500 - 38s - loss: 0.5442 - val_loss: 1.0827 - 38s/epoch - 76ms/step

Epoch 5/100

500/500 - 39s - loss: 0.5401 - val_loss: 1.0847 - 39s/epoch - 77ms/step

Epoch 6/100

500/500 - 39s - loss: 0.5365 - val_loss: 1.0878 - 39s/epoch - 78ms/step

Epoch 7/100

500/500 - 39s - loss: 0.5332 - val_loss: 1.0915 - 39s/epoch - 78ms/step

Epoch 8/100

500/500 - 38s - loss: 0.5302 - val_loss: 1.0956 - 38s/epoch - 75ms/step

Epoch 9/100

500/500 - 39s - loss: 0.5274 - val_loss: 1.0995 - 39s/epoch - 78ms/step

Epoch 10/100

500/500 - 39s - loss: 0.5249 - val_loss: 1.1031 - 39s/epoch - 78ms/step

Epoch 11/100

500/500 - 39s - loss: 0.5225 - val_loss: 1.1066 - 39s/epoch - 77ms/step

Epoch 12/100

500/500 - 39s - loss: 0.5203 - val_loss: 1.1097 - 39s/epoch - 78ms/step


Epoch 13/100

500/500 - 38s - loss: 0.5182 - val_loss: 1.1126 - 38s/epoch - 77ms/step

time: 8min 38s (started: 2025-01-16 02:52:53 +00:00)

```
[28]: def visualize_loss(history, title):  
    loss = history.history["loss"]  
    val_loss = history.history["val_loss"]  
    epochs = range(len(loss))  
    plt.figure()  
    plt.plot(epochs, loss, "b", label="Training loss")  
    plt.plot(epochs, val_loss, "r", label="Validation loss")  
    plt.title(title)  
    plt.xlabel("Epochs")  
    plt.ylabel("Loss")  
    plt.legend()  
    plt.show()
```

```
visualize_loss(history, "Training and Validation LSTM Loss")
```



time: 246 ms (started: 2025-01-16 03:01:31 +00:00)

```
[29]: tf.keras.backend.clear_session()

def build_gru_model(n_steps_in, n_features, n_steps_out):
    model = Sequential([
        GRU(32, activation='relu', input_shape=(n_steps_in, n_features)),
        Dense(n_steps_out)
    ])

    optimizer = tf.keras.optimizers.Adam(1e-4)
    model.compile(optimizer=optimizer, loss=['mse'])
    model.summary()

    return model

# Create and train model
gru_model = build_gru_model(n_steps_in=n_steps_in, n_features=train_X.shape[2],
    ↪n_steps_out=n_steps_out)

early_stopping = tf.keras.callbacks.EarlyStopping(
    monitor='val_loss',
    patience=10,
    mode='min',
    restore_best_weights=True
)

history_gru = gru_model.fit(
    train_X,
    train_y,
    validation_data=(val_X, val_y),
    epochs=100,
    batch_size=32,
    verbose=2,
    shuffle=False,
    callbacks=[early_stopping],
)

gru_model.save('../output/model/gru_model.keras')
```

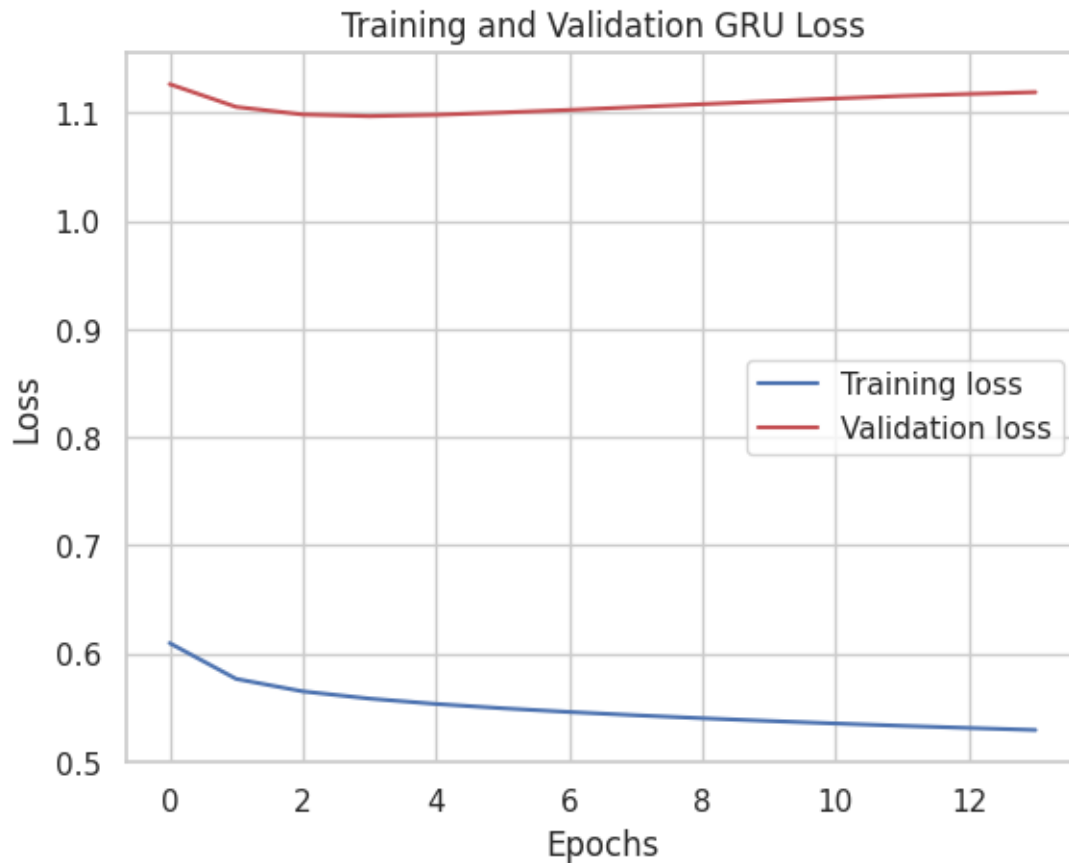
WARNING:tensorflow:Layer gru will not use cuDNN kernels since it doesn't meet the criteria. It will use a generic GPU kernel as fallback when running on GPU.
Model: "sequential"

Layer (type)	Output Shape	Param #
gru (GRU)	(None, 32)	7296

dense (Dense) (None, 1) 33

```
=====
Total params: 7329 (28.63 KB)
Trainable params: 7329 (28.63 KB)
Non-trainable params: 0 (0.00 Byte)
-----
Epoch 1/100
500/500 - 55s - loss: 0.6096 - val_loss: 1.1266 - 55s/epoch - 111ms/step
Epoch 2/100
500/500 - 54s - loss: 0.5765 - val_loss: 1.1055 - 54s/epoch - 108ms/step
Epoch 3/100
500/500 - 53s - loss: 0.5650 - val_loss: 1.0985 - 53s/epoch - 106ms/step
Epoch 4/100
500/500 - 52s - loss: 0.5583 - val_loss: 1.0970 - 52s/epoch - 103ms/step
Epoch 5/100
500/500 - 51s - loss: 0.5533 - val_loss: 1.0982 - 51s/epoch - 103ms/step
Epoch 6/100
500/500 - 52s - loss: 0.5494 - val_loss: 1.1003 - 52s/epoch - 105ms/step
Epoch 7/100
500/500 - 52s - loss: 0.5459 - val_loss: 1.1027 - 52s/epoch - 103ms/step
Epoch 8/100
500/500 - 50s - loss: 0.5429 - val_loss: 1.1054 - 50s/epoch - 100ms/step
Epoch 9/100
500/500 - 52s - loss: 0.5401 - val_loss: 1.1081 - 52s/epoch - 103ms/step
Epoch 10/100
500/500 - 52s - loss: 0.5376 - val_loss: 1.1107 - 52s/epoch - 103ms/step
Epoch 11/100
500/500 - 53s - loss: 0.5353 - val_loss: 1.1133 - 53s/epoch - 106ms/step
Epoch 12/100
500/500 - 54s - loss: 0.5331 - val_loss: 1.1156 - 54s/epoch - 107ms/step
Epoch 13/100
500/500 - 51s - loss: 0.5311 - val_loss: 1.1175 - 51s/epoch - 103ms/step
Epoch 14/100
500/500 - 54s - loss: 0.5293 - val_loss: 1.1191 - 54s/epoch - 109ms/step
time: 12min 21s (started: 2025-01-16 03:01:32 +00:00)
```

```
[30]: visualize_loss(history_gru, "Training and Validation GRU Loss")
```



time: 174 ms (started: 2025-01-16 03:13:53 +00:00)

```
[34]: # Reshape the data to 2D: (samples, features)
train_X_resaped = train_X.reshape(train_X.shape[0], -1) # (15971, 120 * 42)
val_X_resaped = val_X.reshape(val_X.shape[0], -1) # (3903, 120 * 42)

# Initialize the XGBRegressor model
xgb_model = XGBRegressor(
    n_estimators=100, # Number of trees
    learning_rate=0.1, # Learning rate
    max_depth=6, # Maximum depth of trees
    subsample=0.8, # Subsample ratio of training instances
    colsample_bytree=0.8, # Subsample ratio of columns when constructing each
    ↪tree
    random_state=42 # Seed for reproducibility
)

# Train the model
xgb_model.fit(train_X_resaped, train_y.ravel())
```

time: 476 µs (started: 2025-01-16 03:18:05 +00:00)

1.8 Model Evaluation

LSTM Model achieved the lowest MSE and lowest MAE among all models. Demonstrates superior accuracy in predicting the validation data.

XGBRegressor follows closely, suggesting it is also a strong candidate for the task, particularly if computational efficiency is a priority.

```
[35]: from sklearn.metrics import mean_squared_error, mean_absolute_error

# Predictions and evaluation for LSTM model
lstm_val_predictions = lstm_model.predict(val_X)
lstm_mse = mean_squared_error(val_y, lstm_val_predictions)
lstm_mae = mean_absolute_error(val_y, lstm_val_predictions)

# Predictions and evaluation for GRU model
gru_val_predictions = gru_model.predict(val_X)
gru_mse = mean_squared_error(val_y, gru_val_predictions)
gru_mae = mean_absolute_error(val_y, gru_val_predictions)

# Predictions and evaluation for XGB Regressor
xgb_val_predictions = xgb_model.predict(val_X_reshaped)
xgb_mse = mean_squared_error(val_y, xgb_val_predictions)
xgb_mae = mean_absolute_error(val_y, xgb_val_predictions)

# Create a dictionary with the evaluation results
model_comparison = {
    "Model": ["XGBRegressor", "LSTM Model", "GRU Model"],
    "MSE": [xgb_mse, lstm_mse, gru_mse],
    "MAE": [xgb_mae, lstm_mae, gru_mae]
}

# Convert the dictionary to a Pandas DataFrame
result_df = pd.DataFrame(model_comparison)

# Display the DataFrame
print("\nModel Comparison:")
print(result_df)
```

122/122 [=====] - 2s 17ms/step


122/122 [=====] - 2s 16ms/step

Model Comparison:

	Model	MSE	MAE
0	XGBRegressor	1.085550	0.842187

```
1    LSTM Model  1.082546  0.836289
2    GRU Model  1.096962  0.843679
time: 7.75 s (started: 2025-01-16 03:18:11 +00:00)
```

```
[49]: # MSE Bar Chart with Log Scale
plt.figure(figsize=(5, 5))
plt.bar(result_df['Model'], result_df['MSE'], color='blue', alpha=0.7)
plt.title('MSE Comparison Across Models', fontsize=14)
plt.ylabel('MSE Loss', fontsize=12)
plt.xlabel('Models', fontsize=12)
plt.xticks(fontsize=10)
plt.yticks(fontsize=10)
plt.tight_layout()
plt.savefig('../output/visuals/mse_comparison.png', dpi=150,
            bbox_inches='tight')
plt.show()
```



```
time: 214 ms (started: 2025-01-16 03:27:47 +00:00)
```

1.9 Prediction of Global Temperature (Celsius) for the Next Hour

```
[39]: def get_next_hour_prediction(model, last_sequence, scaler, feature_columns):  
    """  
    Get temperature prediction for the next hour using the trained model.  
  
    Args:  
        model: Trained model (LSTM, GRU, or XGBoost)  
        last_sequence: Last n_steps_in days of data (shaped according to model_  
    requirements)  
        scaler: Fitted ColumnTransformer scaler used during training  
        feature_columns: List of feature column names  
  
    Returns:  
        float: Predicted temperature for the next day  
    """  
    # Ensure the input sequence is properly shaped  
    if isinstance(model, XGBRegressor):  
        # For XGBoost, reshape to 2D  
        prediction_input = last_sequence.reshape(1, -1)  
    else:  
        # For LSTM/GRU, keep 3D shape (samples, timesteps, features)  
        prediction_input = last_sequence.reshape(1, last_sequence.shape[0],  
    last_sequence.shape[1])  
  
    # Make prediction  
    prediction = model.predict(prediction_input)  
  
    # Get the RobustScaler from the ColumnTransformer  
    robust_scaler = scaler.named_transformers_['scaler']  
  
    # Create a dummy array with the same shape as the original data  
    dummy_array = np.zeros((1, len(feature_columns)))  
    dummy_array[0, -1] = prediction[0] # Put the prediction in the target_  
    column  
  
    # Inverse transform using the RobustScaler  
    prediction_unscaled = robust_scaler.inverse_transform(dummy_array)[0, -1]  
  
    return prediction_unscaled  
  
# Get the last sequence from your validation data  
last_known_sequence = val_X[-1]  
  
# Get predictions from each model  
# LSTM prediction  
lstm_next_day = get_next_hour_prediction(
```

```

    lstm_model,
    last_known_sequence,
    scaler,
    features
)

# GRU prediction
gru_next_day = get_next_hour_prediction(
    gru_model,
    last_known_sequence,
    scaler,
    features
)

# XGBoost prediction
# Reshape the sequence for XGBoost
last_sequence_reshaped = last_known_sequence.reshape(1, -1)
xgb_next_day = get_next_hour_prediction(
    xgb_model,
    last_sequence_reshaped,
    scaler,
    features
)

# Create an ensemble prediction (simple average)
ensemble_prediction = np.mean([lstm_next_day, gru_next_day, xgb_next_day])

# Print predictions
print("\nNext hour Temperature Predictions:")
print(f"LSTM: {lstm_next_day:.2f}°")
print(f"GRU: {gru_next_day:.2f}°")
print(f"XGBoost: {xgb_next_day:.2f}°")
print(f"Ensemble Average: {ensemble_prediction:.2f}°")

```

1/1 [=====] - 0s 26ms/step

1/1 [=====] - 0s 28ms/step

Next hour Temperature Predictions:

LSTM: 21.86°

GRU: 23.48°

XGBoost: 24.80°

Ensemble Average: 23.38°

time: 135 ms (started: 2025-01-16 03:19:11 +00:00)

/tmp/ipykernel_5464/5228385.py:30: DeprecationWarning: Conversion of an array with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you extract a single element from your array before performing this operation. (Deprecated NumPy 1.25.)

```

dummy_array[0, -1] = prediction[0] # Put the prediction in the target column
/tmp/ipykernel_5464/5228385.py:30: DeprecationWarning: Conversion of an array
with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you
extract a single element from your array before performing this operation.
(Deprecated NumPy 1.25.)
dummy_array[0, -1] = prediction[0] # Put the prediction in the target column

```

```

[40]: from datetime import datetime, timedelta

def visualize_hourly_predictions(models, last_sequence, scaler, features,
    ↪history_hours=48):
    """
    Visualize the last N hours of actual data and predicted next hour
    ↪temperature.

    Args:
        models: Dictionary of models {'name': model_object}
        last_sequence: Last sequence of data used for prediction
        scaler: Fitted ColumnTransformer scaler
        features: List of feature column names
        history_hours: Number of past hours to display (default: 48)
    """
    # Get the RobustScaler from the ColumnTransformer
    robust_scaler = scaler.named_transformers_['scaler']

    # Create figure and axis with larger size
    plt.figure(figsize=(12, 6))

    # Plot historical data
    historical_data = last_sequence[-history_hours:, -1] # Get last column
    ↪(temperature)
    historical_times = np.arange(-history_hours, 0)

    # Inverse transform historical data
    historical_data_reshaped = np.zeros((len(historical_data), len(features)))
    historical_data_reshaped[:, -1] = historical_data
    historical_data_unscaled = robust_scaler.
    ↪inverse_transform(historical_data_reshaped)[:, -1]

    # Plot historical data
    plt.plot(historical_times, historical_data_unscaled,
             label='Historical', color='gray', linewidth=2)

    # Get and plot predictions for each model
    colors = ['blue', 'green', 'red']
    predictions = []

```



```

for (name, model), color in zip(models.items(), colors):
    # Reshape input based on model type
    if isinstance(model, XGBRegressor):
        prediction_input = last_sequence.reshape(1, -1)
    else:
        prediction_input = last_sequence.reshape(1, last_sequence.shape[0],
↪last_sequence.shape[1])

    # Get prediction
    pred = model.predict(prediction_input)[0]

    # Inverse transform prediction
    dummy_array = np.zeros((1, len(features)))
    dummy_array[0, -1] = pred
    pred_unscaled = robust_scaler.inverse_transform(dummy_array)[0, -1]
    predictions.append(pred_unscaled)

    # Plot prediction point
    plt.plot([0, 1], [historical_data_unscaled[-1], pred_unscaled],
             color=color, linestyle='--', label=f'{name} Prediction')
    plt.scatter([1], [pred_unscaled], color=color, s=100)

    # Calculate and plot ensemble prediction
    ensemble_pred = np.mean(predictions)
    plt.plot([0, 1], [historical_data_unscaled[-1], ensemble_pred],
             color='purple', linestyle='--', label='Ensemble Prediction')
    plt.scatter([1], [ensemble_pred], color='purple', s=100)

    # Customize the plot
    plt.grid(True, linestyle='--', alpha=0.7)
    plt.title('Hourly Temperature Prediction', fontsize=14, pad=20)
    plt.xlabel('Hours (Past → Future)', fontsize=12)
    plt.ylabel('Temperature (°)', fontsize=12)

    # Set x-axis ticks
    plt.xticks(np.arange(-history_hours, 2, 6))

    # Add legend
    plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left')

    # Adjust layout to prevent label cutoff
    plt.tight_layout()

    return plt.gcf()

# Create dictionary of models
models = {

```

```

'LSTM': lstm_model,
'GRU': gru_model,
'XGBoost': xgb_model
}

# Visualize predictions
fig = visualize_hourly_predictions(
    models=models,
    last_sequence=last_known_sequence,
    scaler=scaler,
    features=features,
    history_hours=48 # Show last 48 hours
)

plt.savefig('../output/visuals/prediction_next_hour.png', dpi=300,
            bbox_inches='tight')
plt.show()

```

```

1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 31ms/step


```

/tmp/ipykernel_5464/1936604855.py:49: DeprecationWarning: Conversion of an array with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you extract a single element from your array before performing this operation.
(Deprecated NumPy 1.25.)

```
dummy_array[0, -1] = pred
```

/tmp/ipykernel_5464/1936604855.py:49: DeprecationWarning: Conversion of an array with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you extract a single element from your array before performing this operation.
(Deprecated NumPy 1.25.)

```
dummy_array[0, -1] = pred
```



time: 876 ms (started: 2025-01-16 03:19:14 +00:00)

[]: